

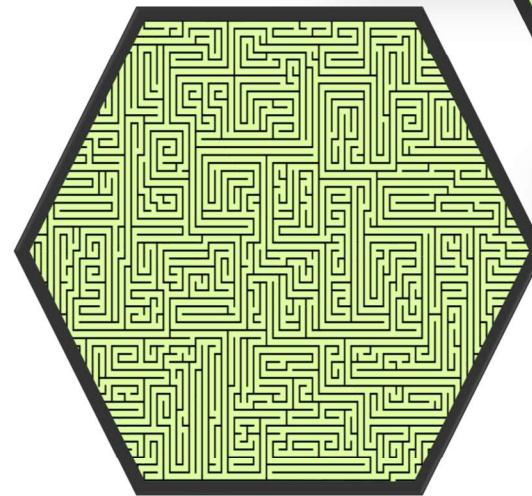
Lecture 04

# OPTIMISATION TECHNIQUES

Tilo Burghardt | [tilo@cs.bris.ac.uk](mailto:tilo@cs.bris.ac.uk)

# Agenda for Lecture 4

- Recap Backpropagation
- Optimisation Techniques



# RECAP: BACKPROPAGATION ALGORITHM



# Recap: Backpropagation Algorithm

**initialise** all weights  $w_{ij}^l$  randomly

**for**  $t=0, 1, 2, \dots$  **do**

**pick** next training sample  $([f_1^0, f_2^0, \dots], [f_1^*, f_2^*, \dots])$

**FORWARD PASS:** compute all  $s_j^l = \sum_{i=1}^{d(l-1)} w_{ij}^l f_i^{l-1}$  and  $f_j^l = g_j^l(s_j^l)$

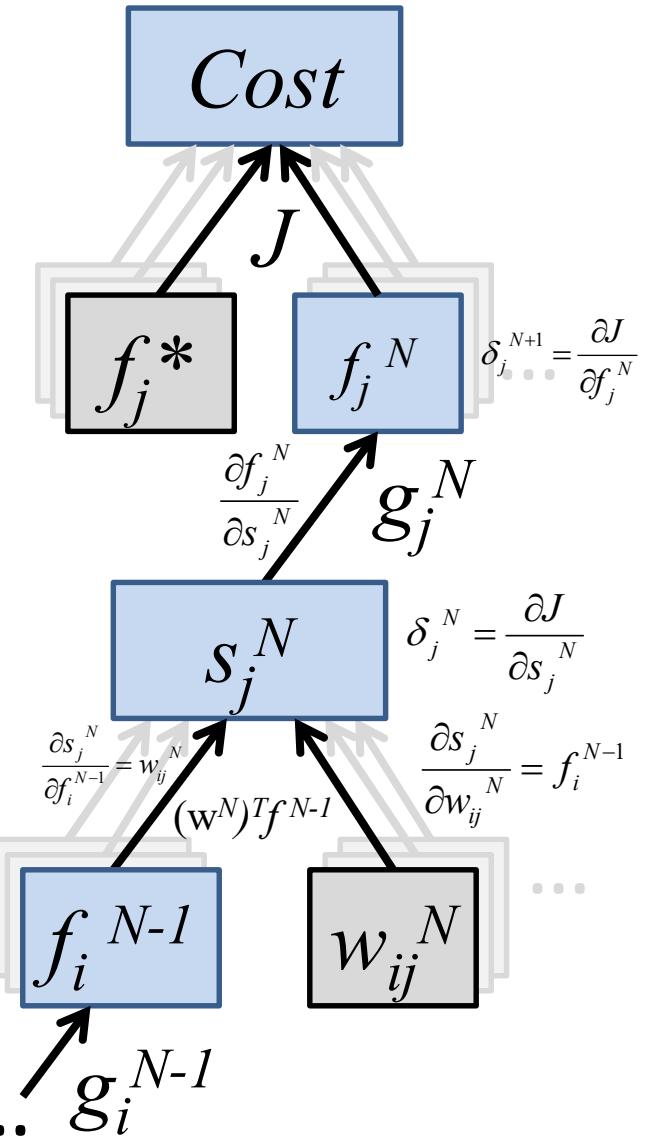
**compute** top deltas  $\delta_j^N = g_j^N(s_j^N) \cdot \partial J / \partial f_j^N$

**BACKWARD PASS:** compute all  $\delta_i^{l-1} = g_i^{l-1}(s_i^{l-1}) \sum_{j=1}^{d(l)} w_{ij}^l \delta_j^l$

**update** weights  $w_{ij}^l \leftarrow w_{ij}^l - \eta f_i^{l-1} \delta_j^l$

**check** if stopping criteria are met to break loop

**return** final weights  $w_{ij}^l$



# SGD



# (Online) Backpropagation so far: Notational Compaction

initialise all weights  $w_{ij}^l$  randomly

**for**  $t=0, 1, 2, \dots$  **do**

**pick** next training sample  $([f_1^0, f_2^0, \dots], [f_1^*, f_2^*, \dots])$

**FORWARD PASS:** compute all  $s_j^l = \sum_{i=1}^{d(l-1)} w_{ij}^l f_i^{l-1}$  and  $f_j^l = g_j^l(s_j^l)$

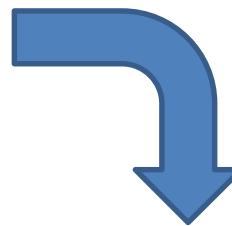
**compute** top deltas  $\delta_j^N = g_j^N(s_j^N) \cdot \partial J / \partial f_j^N$

**BACKWARD PASS:** compute all  $\delta_i^{l-1} = g_i^{l-1}(s_i^{l-1}) \sum_{j=1}^{d(l)} w_{ij}^l \delta_j^l$

**update** weights  $w_{ij}^l \leftarrow w_{ij}^l - \eta f_i^{l-1} \delta_j^l$

**check** if stopping criteria are met to break loop

**return** final weights  $w_{ij}^l$



initialise all weights  $W$  randomly

**for**  $t=0, 1, 2, \dots$  **do**

**pick** next training sample  $(x, f^*)$

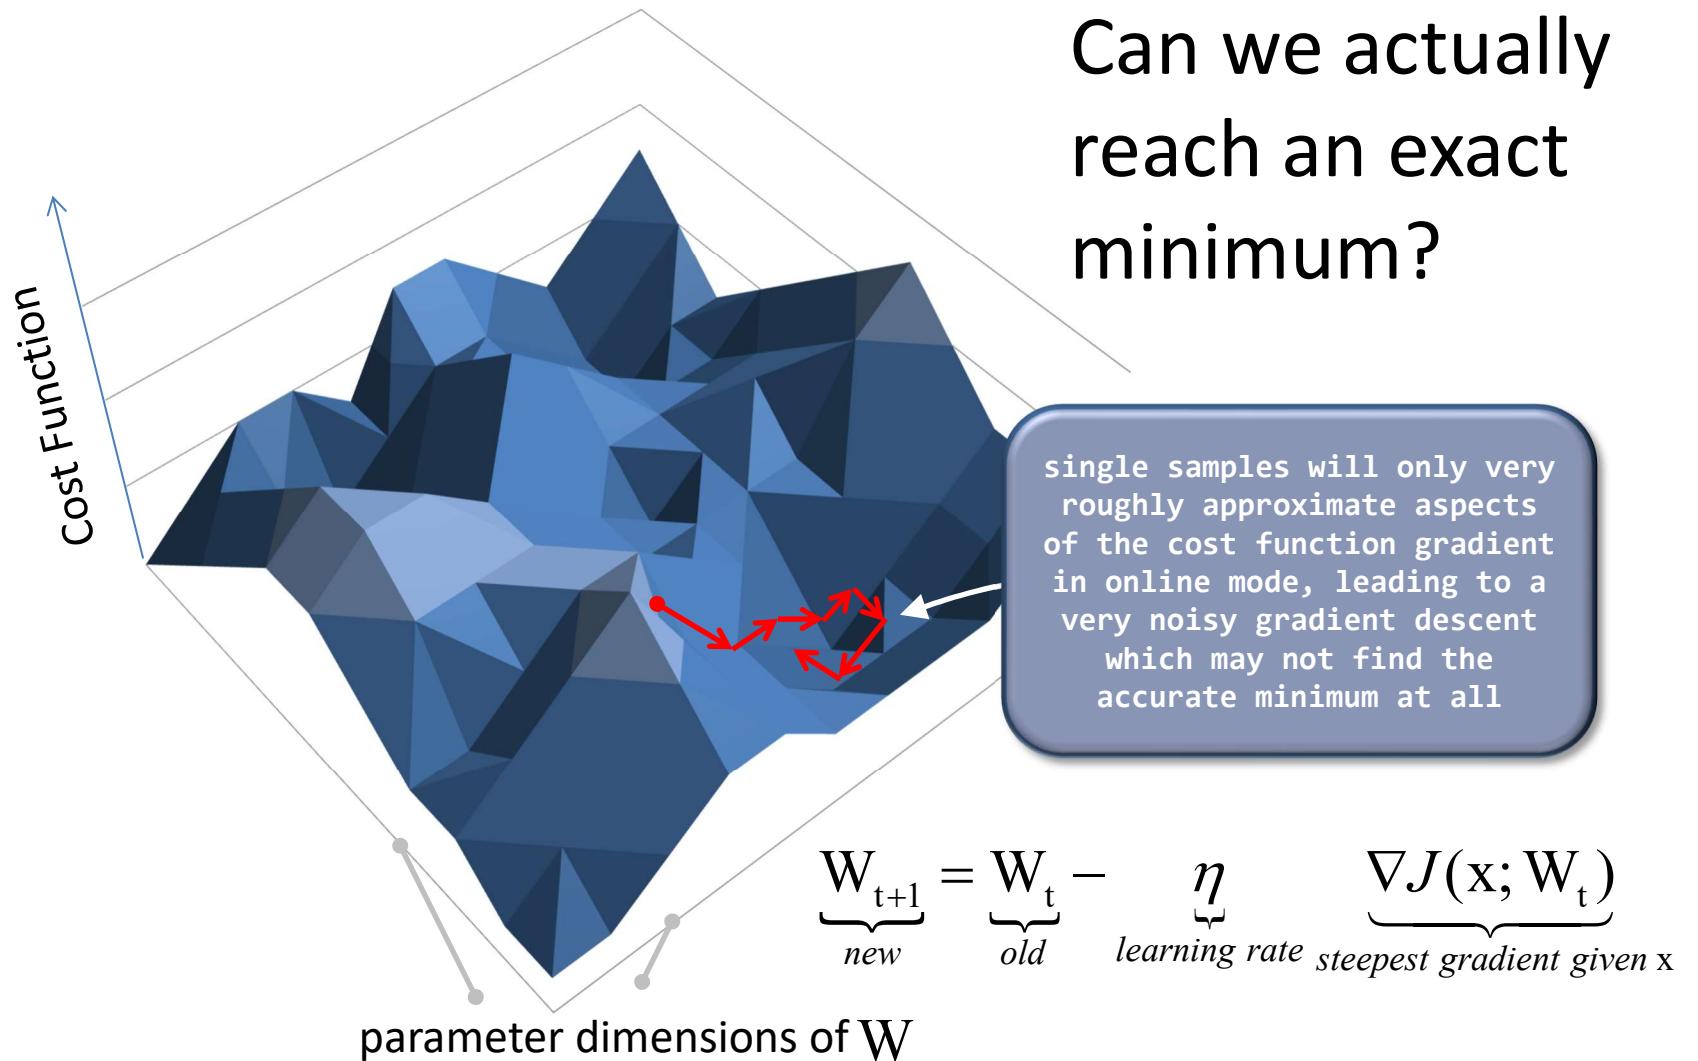
**FORWARD-BACKWARD PASS:** compute  $\nabla J$

**update** weights  $W \leftarrow W - \eta \nabla J$

**check** if stopping criteria are met to break loop

**return** final weights  $W$

# Noisy Gradient Descent due to Online Sampling



# Online, Deterministic and Stochastic Training

Can we actually  
reach an exact  
minimum?

(MINIBATCH)  
STOCHASTIC  
GRADIENT  
DESCENT

```
initialise all weights W randomly
for t=0, 1, 2, ... do
    pick a small subset of training samples (X, F*)
    FORWARD-BACKWARD PASS: compute ∇J
    update weights W ← W - η∇J
    check if stopping criteria are met to break loop
return final weights W
```

$$\nabla J = \frac{1}{|X|} \nabla_W \sum_j L(f(x_j, W), f^*)$$

ONLINE  
GRADIENT  
DESCENT

```
initialise all weights W randomly
for t=0, 1, 2, ... do
    pick next training sample (x, f*)
    FORWARD-BACKWARD PASS: compute ∇J
    update weights W ← W - η∇J
    check if stopping criteria are met to break loop
return final weights W
```

initialise all weights W randomly

for t=0, 1, 2, ... do

use all training samples (X, F\*)

FORWARD-BACKWARD PASS: compute ∇J

update weights W ← W - η∇J

check if stopping criteria are met to break loop

return final weights W

**DETERMINISTIC GRADIENT DESCENT**

given small enough learning rate DGD will make progress to true local minimum, but at high computational cost!

# Practical Solution: SGD with ‘Simulated Annealing’

**initialise** all weights  $\mathbf{W}$  randomly

**for**  $k=0, 1, 2, \dots \tau$  **do**

$$\eta_k = (1 - \frac{k}{\tau})\eta_0 + \frac{k}{\tau}\eta_\tau$$

introduction of a changing learning rate  $\eta_k$  decreasing over  $\tau+1$  steps by blending from a starting learning rate  $\eta_0$  towards a final learning rate  $\eta_\tau$ .

**for**  $t=0, 1, 2, \dots$  **do**

**pick** a small subset of training samples  $(\mathbf{X}, \mathbf{F}^*)$

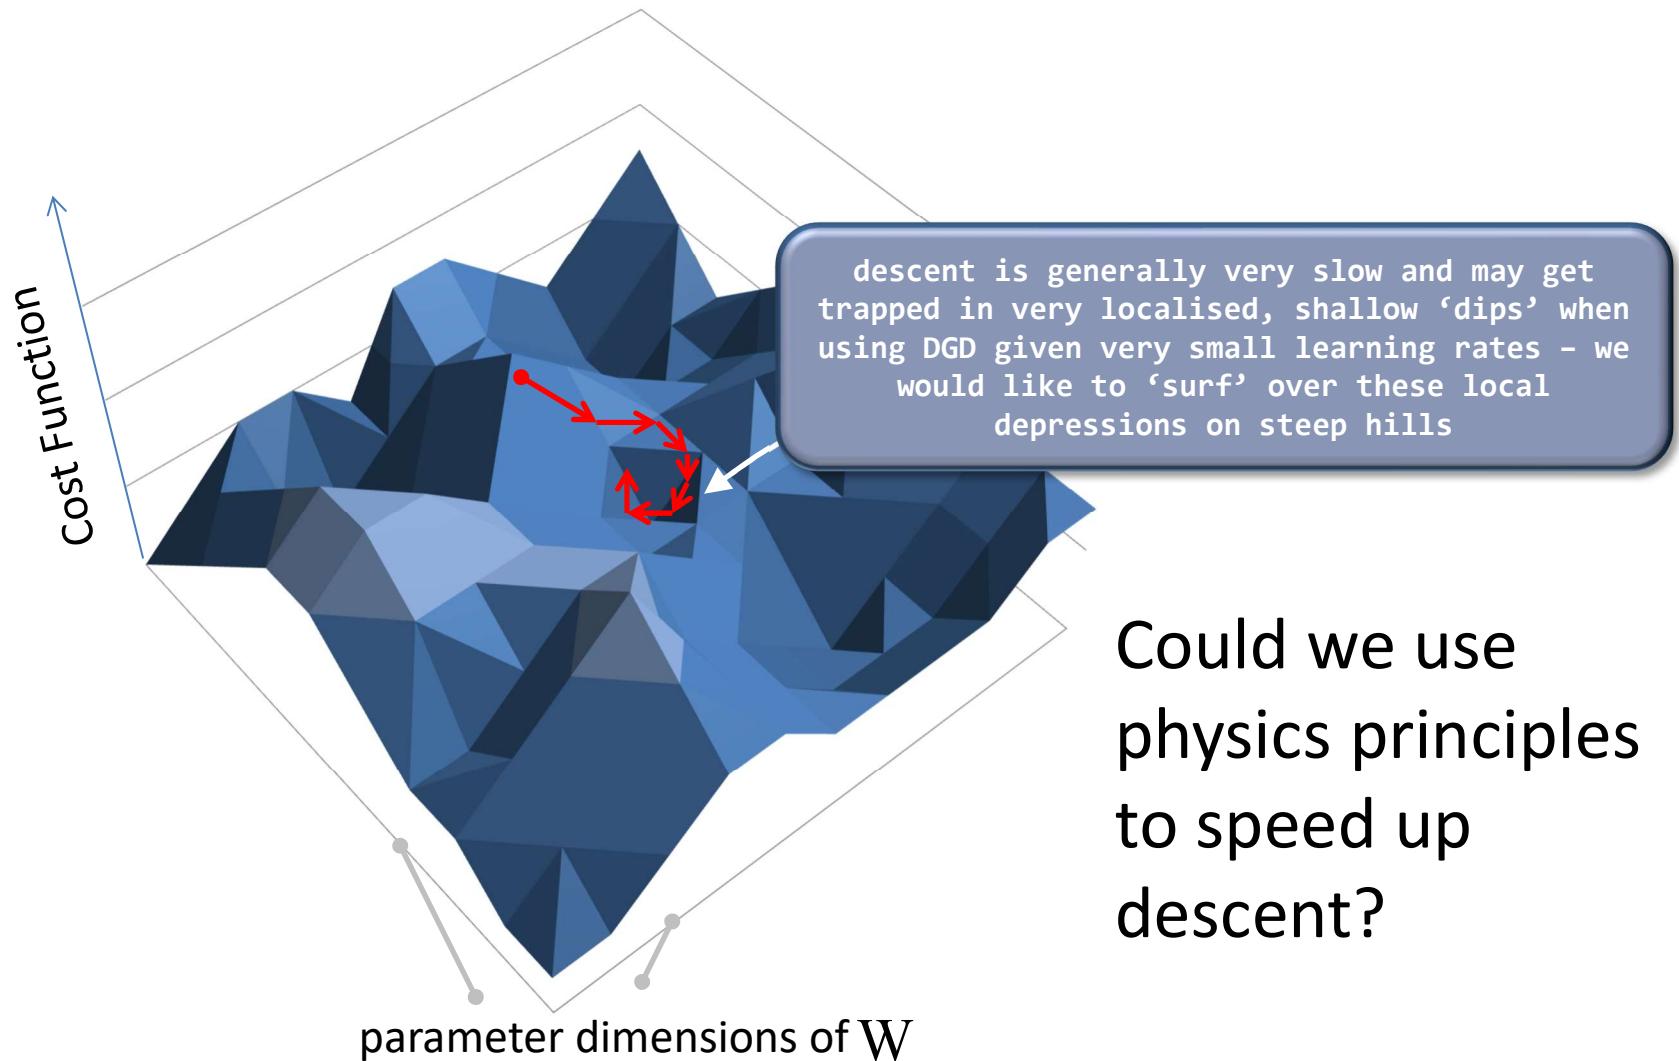
**FORWARD-BACKWARD PASS:** compute  $\nabla J$

**update** weights  $\mathbf{W} \leftarrow \mathbf{W} - \eta_k \nabla J$

...

**return** final weights  $\mathbf{W}$

# Slow Descent and Local ‘Dips’ of Cost Function



# MOMENTUM



# Speeding up Learning via Momentum

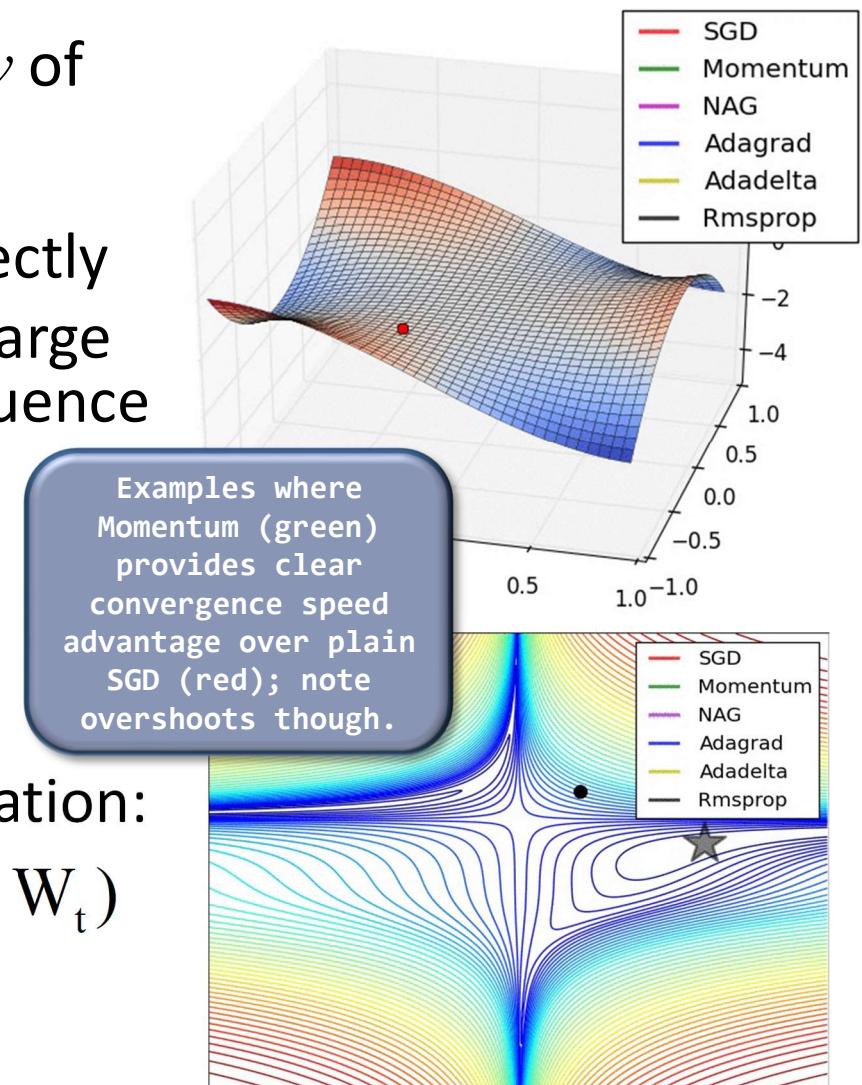
- **Idea:** introduce a velocity term  $v$  of ‘current descent speed’ and use current gradient to change this velocity rather than weights directly
- step sizes now depend on how large and how aligned a previous sequence of gradients has been
- formally, we change the update equations for weights from:

$$\mathbf{W}_{t+1} = \mathbf{W}_t - \eta \nabla J(\mathbf{X}; \mathbf{W}_t)$$

by introducing velocity accumulation:

$$v_{t+1} = \underbrace{\alpha}_{\text{momentum parameter}} v_t - \eta \nabla J(\mathbf{X}; \mathbf{W}_t)$$

$$\mathbf{W}_{t+1} = \mathbf{W}_t + \underbrace{v_{t+1}}_{\text{momentum}}$$



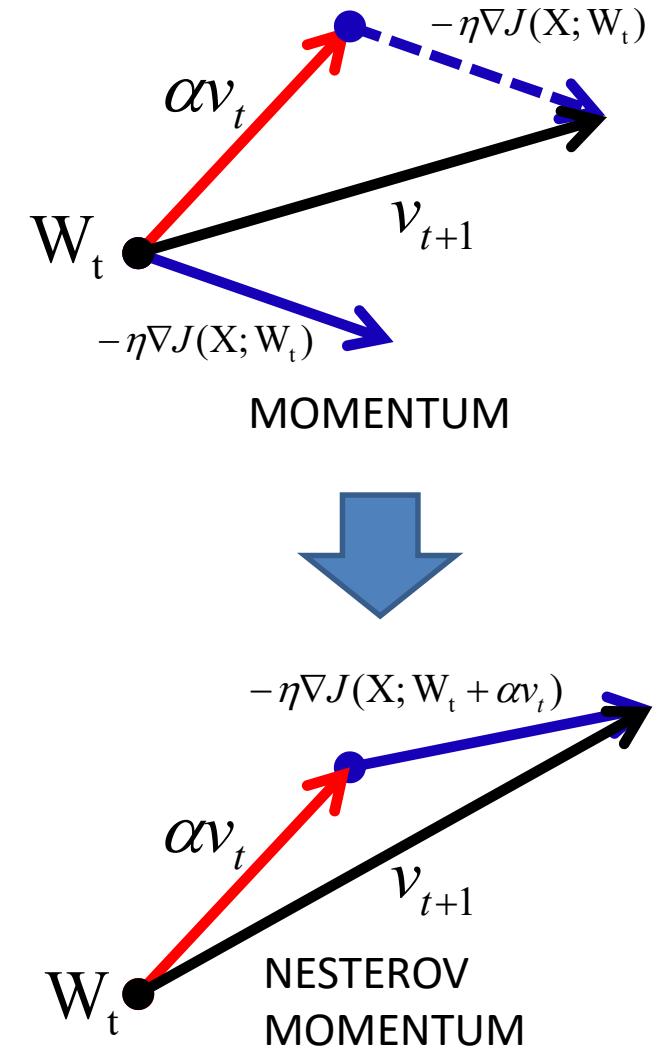
animation sources:  
Alec Radford

# Nesterov Accelerated Gradient (NAG)

- **Idea:** don't calculate gradient at current position since momentum will carry us forward to another position anyway – take (lookahead) gradient at target
- can be seen as adding a 'correction term' to the standard method of momentum
- consistently works slightly better than standard momentum in practice
- weights are now updated as follows:

$$\nu_{t+1} = \alpha v_t - \eta \nabla J(X; \underbrace{W_t + \alpha v_t}_{\text{preview location}})$$

$$W_{t+1} = W_t + \nu_{t+1}$$



- however, still very slow progress on shallow plateau regions

# NEWTON'S METHOD



# Newton's Method (2<sup>nd</sup> Order)

- **Idea:** let curvature rescale the gradient – multiplying the gradient by the inverse Hessian leads to an optimization that takes aggressive steps in directions of shallow curvature and shorter steps in directions of steep curvature
- great advantage: no extra learning rate or hyperparameters needed
- however, computing and inverting the Hessian is very expensive and space consuming (Hessian  $\mathbf{H}$  has square size w.r.t. to number of weights!):

$$\mathbf{W}_{t+1} = \mathbf{W}_t - \mathbf{H}(J(\mathbf{X}; \mathbf{W}_t))^{-1} \nabla J(\mathbf{X}; \mathbf{W}_t)$$

- yet, *Newton's method* without modifications has a critical shortcoming: it is attracted to Saddle points... (see also “Hessian-free” 2<sup>nd</sup>-order methods)

RECAP: HESSIAN MATRIX

$$\mathbf{H}(J) = \begin{bmatrix} \cdots & \cdots & \cdots \\ \cdots & \frac{\partial^2 J}{\partial w_i \partial w_j} & \cdots \\ \cdots & \cdots & \cdots \end{bmatrix}$$

# Saddle Points as Critical Points

- There are various point categories of the objective function where the gradient is zero:
  - **Minima** (all Eigenvalues of Hessian positive),
  - **Maxima** (all Eigenvalues of Hessian negative),
  - **Saddle points** (both positive and negative Eigenvalues of Hessian)

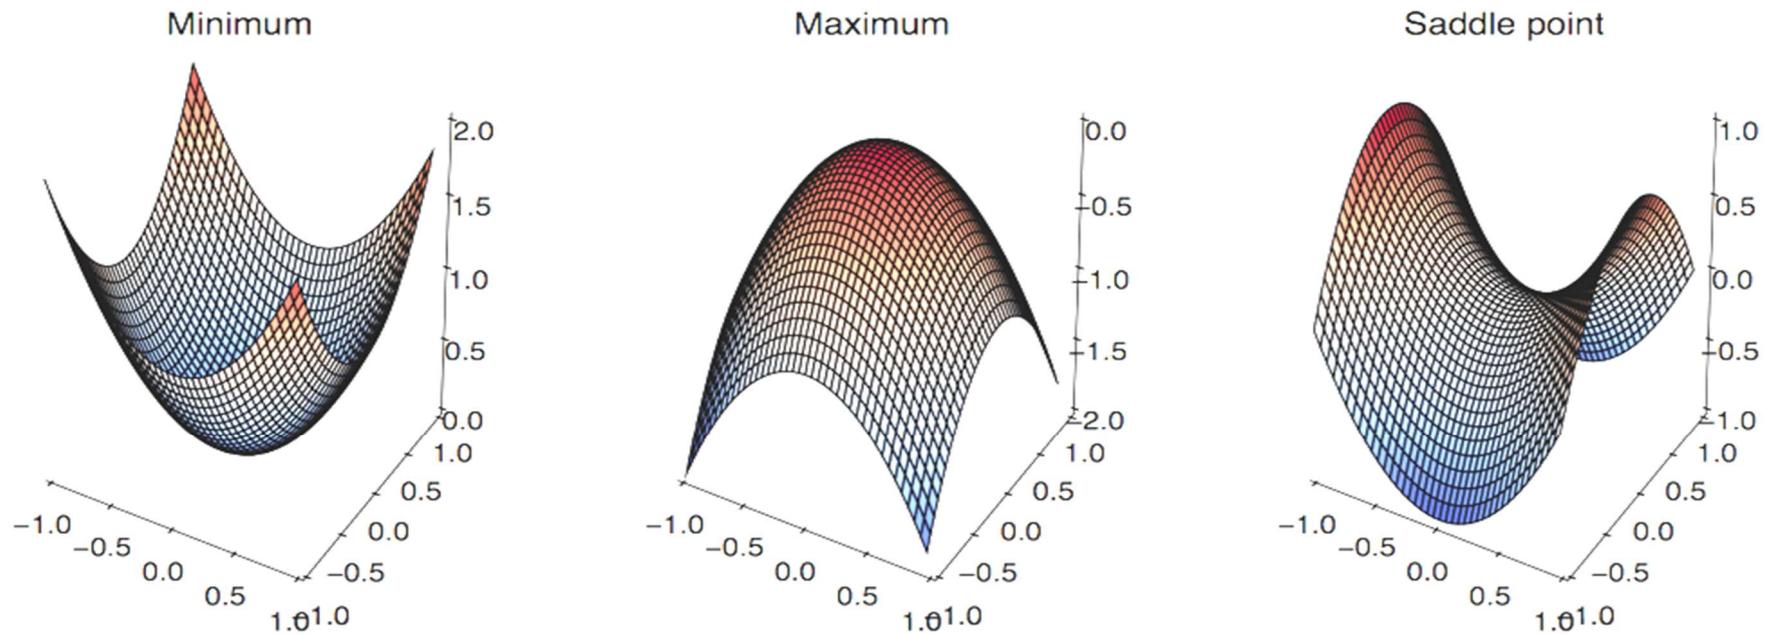
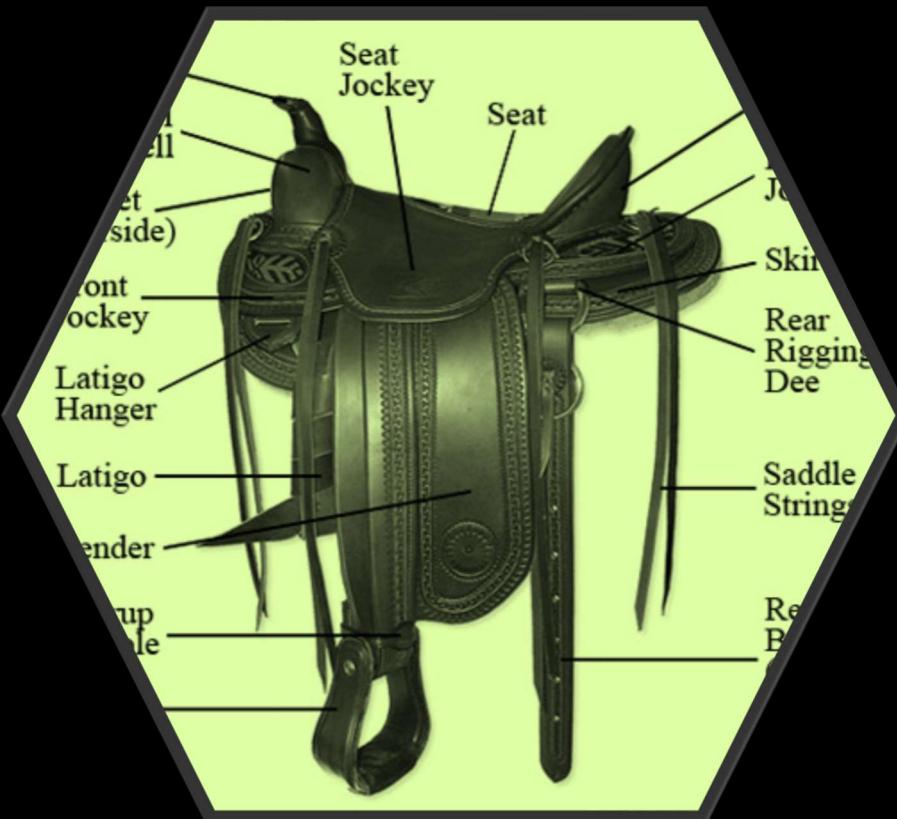


image adjusted from source: Ian Goodfellow 2015

# SADDLE POINT CONSIDERATIONS



# Are there more Saddle points or local Minima?

- for an arbitrary problem, assume sign of Hessian Eigenvalues is random:
  - exponentially less likely to get ‘all positive’ (i.e. being a Minimum) with higher and higher parameter dimensions
- Random Matrix Theory provides further insight:
  - the lower  $J$  is, the more likely to find positive Eigenvalues
- neural nets without non-linearities have global minima connected via a single manifold and many Saddle points (Saxe et al, 2013)

## GOOD NEWS:

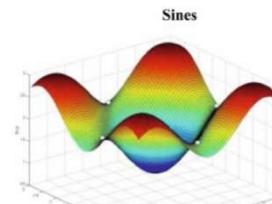
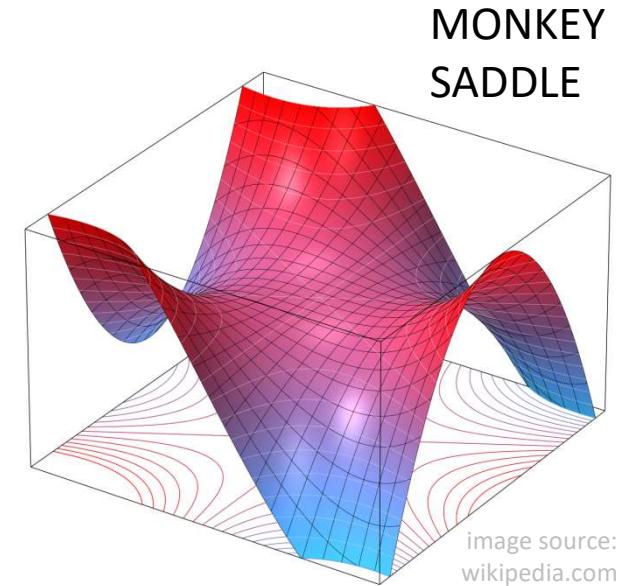
- Most critical points with higher cost  $J$  should be Saddle points and they offer a chance to escape from them particularly via symmetry-breaking descent-methods!
- Most local minima should therefore have a low cost  $J$  associated with them and *may* be reachable via descent!

# High Number of Saddle points

- experiments and theoretical arguments (Dauphin et al 2014, Choromanska et al 2015) provide some support that neural nets have indeed as many Saddle points as Random Matrix Theory proposes
- in fact, the number of Saddle points may increase exponentially with the dimensionality of the function

## NOT SO GREAT NEWS:

- *Newton's method* will work poorly (since being attracted to Saddle points) with a high chance of getting stuck
- however, idea of a function-adaptive learning rate seems valuable



SADDLE POINT EXAMPLES  
(white)

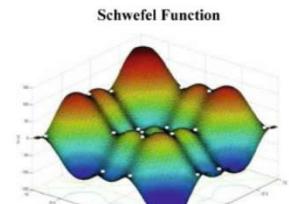
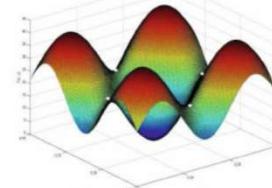


image source:  
Chen et al. 2012

# PER-WEIGHT ADAPTIVE GRADIENTS



# Adagrad (adaptive gradient) (*Duchi et al. 2011*)

- **Idea:** keep track of per-weight learning rates to force evenly spread learning speeds – weights that are associated with high gradients have their effective rate of learning decreased, whilst weights that have infrequent or particularly small updates have their rates increased.
- such ‘monotonic learning’ may help with issues including breaking of symmetries and slow progress in particular dimensions
- update now uses a  $W_t$ -sized accumulator  $A$ :

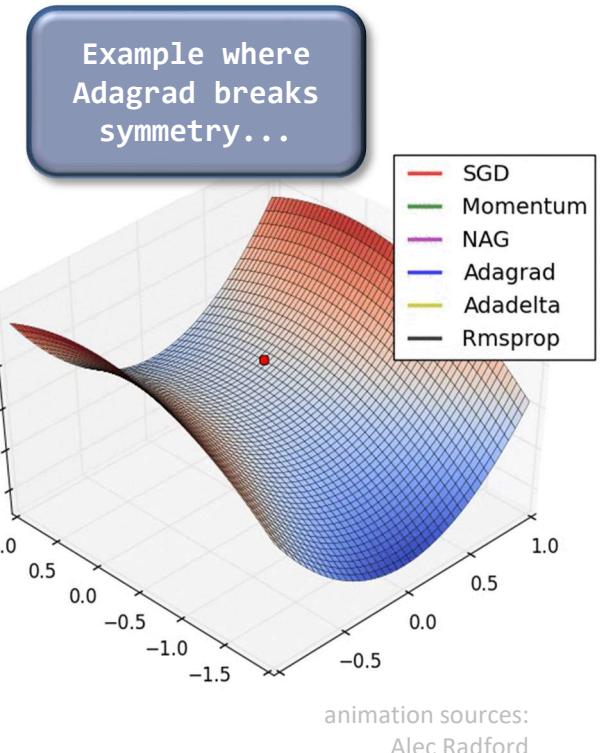
$$A_{t+1} = A_t + (\nabla J(X; W_t))^2$$

element-by-element squaring

$$W_{t+1} = W_t - \eta \frac{\nabla J(X; W_t)}{(\sqrt{A_{t+1}} + \epsilon)}$$

avoiding division by zero

- however, this ‘monotonic learning’ is a very aggressive approach and lacks the possibility of late adjustments...learning usually stops too early...



Concept: element-wise dampening of historically highly active gradient components  
( $A$  being large) and amplification of slowly changing gradient components  
( $A$  being small)

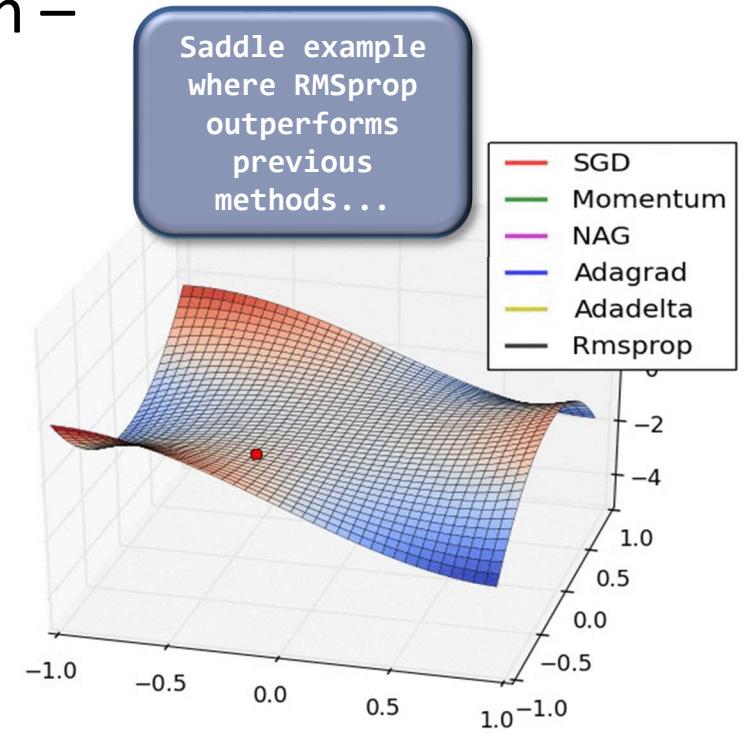
# RMSprop (*Hinton “L.6 S.29”*)

- **Idea:** root-mean-square propagation – combat the aggressive reduction in Adagrad’s learning speed by propagation of a smooth running average
- update equations now introduce a smoothing parameter  $\beta$ :

$$A_{t+1} = \beta A_t + (1 - \beta)(\nabla J(X; W_t))^2$$

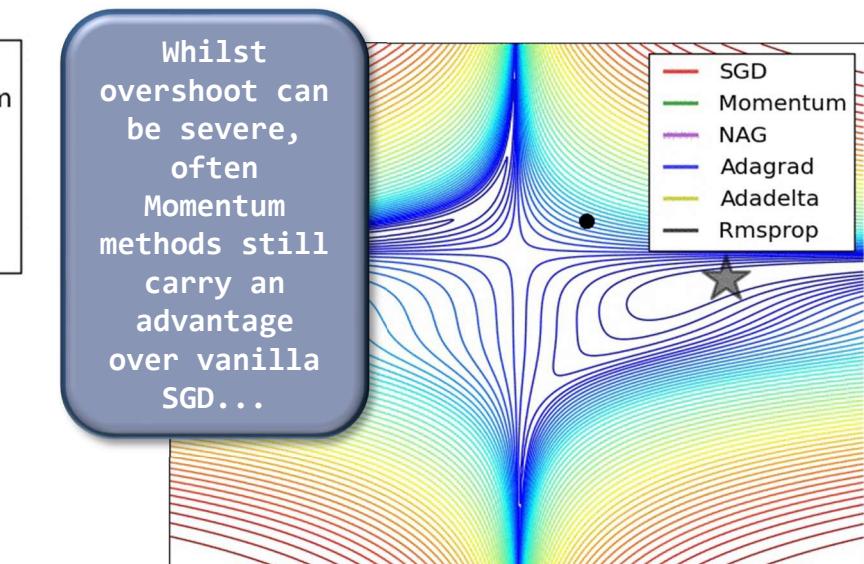
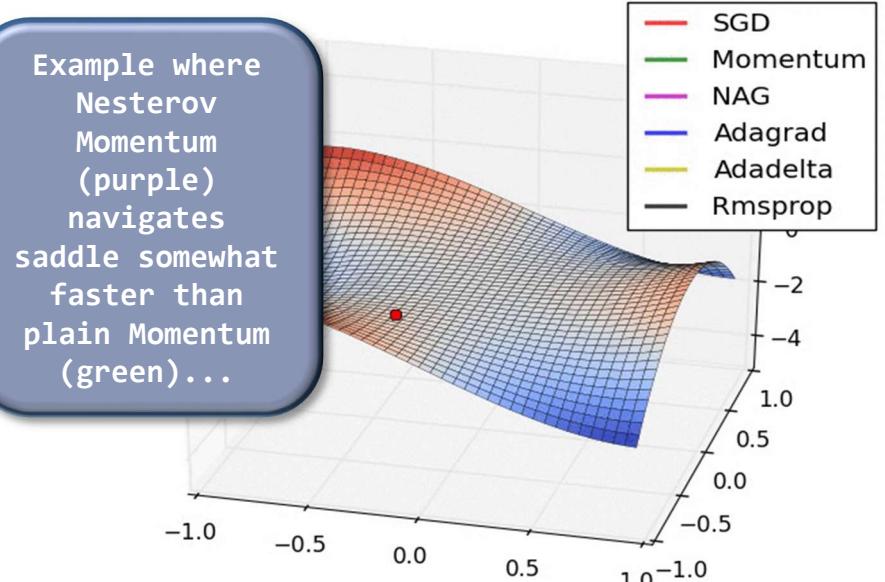
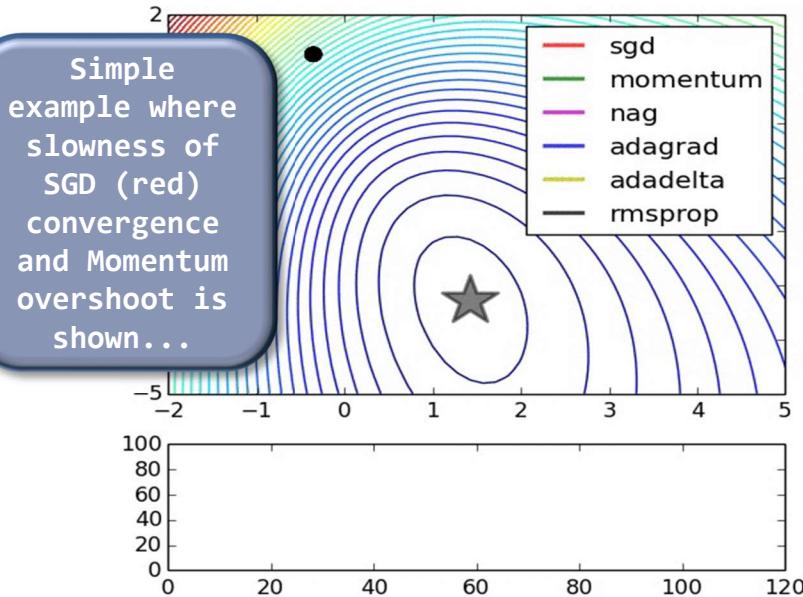
$$W_{t+1} = W_t - \eta \frac{\nabla J(X; W_t)}{(\sqrt{A_{t+1}} + \epsilon)}$$

- just adding standard momentum does not help much in improving performance further (*see Hinton*)
- however, further smoothing and correction operations can be applied...



Concept: element-wise dampening of recently highly active gradient components (A being large) and amplification of slowly changing gradient components (A being small)

# Some Observations in Convergence Visualisations



animation sources:  
Alec Radford

# ADAM



# Adam<sub>(adaptive moment estimation)</sub> (*Kingma & Ba 2014*)

- **Idea 1:** smooth RMSprop's usually 'noisy' incoming gradient (beyond the effect of mini-batching) using a new parameter  $\alpha$ :

$$\begin{aligned} G_{t+1} &= \alpha G_t + (1 - \alpha) \nabla J(X; W_t) \\ A_{t+1} &= \beta A_t + (1 - \beta) (\nabla J(X; W_t))^2 \\ W_{t+1} &= W_t - \eta \frac{G_{t+1}}{(\sqrt{A_{t+1}} + \varepsilon)} \end{aligned}$$

- **Idea 2:** correct for the impact of bias introduced by 'initialising' the two smoothed measures – i.e. starting with  $t=1$  'fade-in' the smoothing effect exponentially by introducing  $\bar{G}$  and  $\bar{A}$  :

$$\begin{aligned} G_{t+1} &= \alpha G_t + (1 - \alpha) \nabla J(X; W_t) \\ \bar{G} &= G_{t+1} / (1 - \alpha^t) \\ A_{t+1} &= \beta A_t + (1 - \beta) (\nabla J(X; W_t))^2 \\ \bar{A} &= A_{t+1} / (1 - \beta^t) \\ W_{t+1} &= W_t - \eta \frac{\bar{G}}{(\sqrt{\bar{A}} + \varepsilon)} \end{aligned}$$

# Summary Adam (adaptive moment estimation)

fading of gradient  
→ force fast start-up

$$\bar{G}_{t+1} = \alpha G_t + (1 - \alpha) \nabla J(X; W_t)$$

smoothing of gradients

fading of per-weight magnitudes

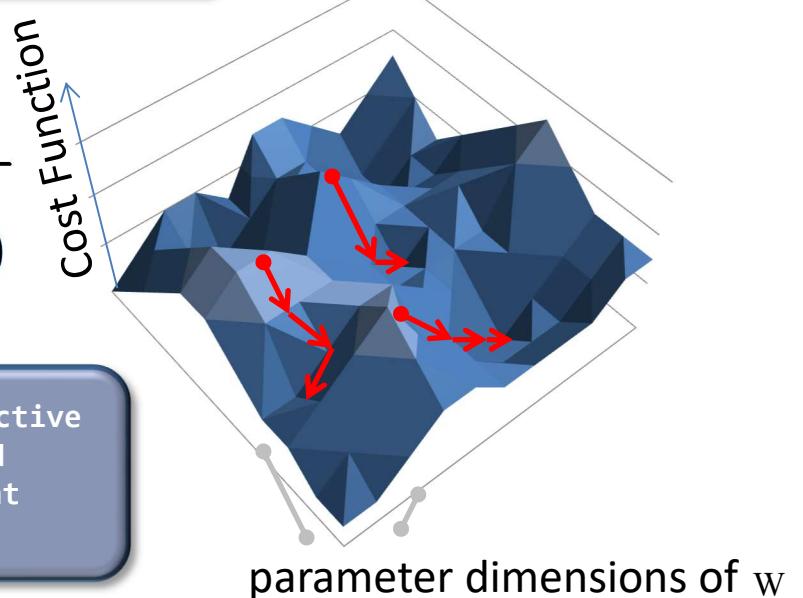
$$\bar{A}_{t+1} = \beta A_t + (1 - \beta) (\nabla J(X; W_t))^2$$

RMS propagation

element-wise squaring

$$W_{t+1} = W_t - \eta \frac{\bar{G}}{(\sqrt{\bar{A}} + \epsilon)}$$

element-wise dampening of recently highly active gradient components ( $A$  being large) and amplification of slowly changing gradient components ( $A$  being small)



# Right, can we train deep networks now? – Maybe...

- Why is applying Adam to ReLU-based networks not a guarantee for successful deep learning then?
    - We have introduced new parameters  $\alpha$ ,  $\beta$ ,  $\varepsilon$ ... : how to set these so-called ‘hyper-parameters’?
    - Even our mini-batch size has not been discussed...
    - We have not talked about network initialisation – this matters a lot and can change results drastically if done wrong.
    - Overfitting is likely to occur in deep networks as in any learning system: regularisation techniques are critical to achieve good generalisation beyond the training data available!
    - Number of parameters explodes in deep networks; we may need to share them or reuse the entire net (e.g. CNNs/RNNs).
    - The simple loss functions discussed so far need extending to provide better results for common tasks such as classification.
    - The data we deal with is part of the training process – we have not talked about data at all so far...
- Yet, applying deep learning and achieving top-end results often involves a lot of parameter tuning, testing and trial-and-error of various designs and techniques available, and performance is critically dependent on the quality of training data and also the GPU-sizes which limit network designs – it is still as much an ‘engineering process’ as it is a science...

# Next: COST FUNCTIONS, REGULARISATION AND DEPTH

- Key Loss Functions
- L1 and L2 Weight Decay
- Dropout and Noise
- Data Augmentation
- Why deep is advantageous...
- Scalability Considerations...

